Dear Students,

For the project phase of the course, NVIDIA is providing us with GPU compute through Brev (https://brev.nvidia.com). This email contains everything you need to get set up, plus the rules we expect every team to follow.

First action: nominate a Team Captain

Each team nominates one Team Captain and has them **reply to all people CC'd in this email** (from their ETH address). The captain is responsible for:

- Creating the team's Brev organization
- Redeeming the coupon code
- Inviting teammates
- Sharing GPU instances (SSH / Jupyter) with the team
- Keeping an eye on the team's usage

We can't issue the coupon until we know who your captain is.

How the credits are structured

- Credits are allocated per team, not per student.
- Each team's baseline allocation is $200 of Brev credit, delivered as a single coupon code redeemed by the captain.

Captain: redeem the coupon

After you reply nominating yourself as captain, we will email the coupon code directly to you. Then:

1. Go to https://brev.nvidia.com and sign up (please use your ETH email).
2. Open the **Billing** tab, scroll to **Redeem Code**, paste the code, click **Redeem**.
3. Open the **Team** tab, click **Generate Invite Link**, and share it with your teammates.

Teammates: join the org

Each teammate clicks the invite link and signs in to Brev with their ETH email.

Sharing a GPU instance with your team

Once the captain has spun up an instance:

- **SSH:** open the instance in the GPU tab → scroll to **Share SSH Access** → add your teammates' usernames. The instance must be running for this to work.
- **Jupyter:** under **Using secure links**, click **Edit access**, add the teammate's email/username, click **Share**. The container must be fully built before this works.

Hardware you should expect

- H100 is the target GPU. A100 is an acceptable fallback when H100 capacity is tight.
- Size your instance to the job. Don't default to the largest node you can find: a smaller instance is usually enough, and your $200 lasts proportionally longer.

- (recommendation) While prototyping, stick to a single GPU until your code works. Most bugs show up just as clearly on one GPU as on eight (and debugging on 8×H100s burns credits fast). Scale up once the pipeline runs end-to-end.
- Brev pulls instances from multiple cloud providers, and availability changes dynamically based on load.
- (recommendation) Pick one provider and stick with it for your project. Different providers have subtly different environments, and swapping mid-project is a reliable way to introduce bugs that are painful to debug. When you first launch a working setup, write down which provider you used.

Using credits responsibly

- Stop or terminate instances the moment you stop using them. Idle GPUs burn credits very fast.
- Pause/resume is provider-dependent. Some providers do not support stop/start. On those, terminating means you will need to reupload your data. Data is tied to the specific instance.
- Going away for a while (e.g. the weekend)? `scp` your data out and terminate the instance instead of leaving it running.
- Negative-balance instances are deleted after 96 hours. Monitor your balance in the Billing tab. If you overshoot and the instance gets deleted, your data goes with it.

Required: add us to your Brev org

Every team must add the following two TAs to their Brev team after setup:

- Zador Pataki (patakiz@ethz.ch)
- Nicole Damblon (ndamblon@ethz.ch)

Add us via the Team tab using the same invite link you send to your teammates. We use this to monitor usage across all 50 teams and to help debug problems.

Getting help

- **Brev platform issues** (SSH, billing, provider problems, how to `scp` data out, etc.): brev-support@nvidia.com
- **Brev docs and in-product AI search:** https://docs.nvidia.com/brev/latest (also accessible from the Documentation tab inside the Brev console). Please read the docs before asking us or Brev support.
- **Course-specific questions, coupon requests:** Zador (patakiz@ethz.ch), Nicole (ndamblon@ethz.ch)

Good luck with your projects. We're excited to see what you build.

Best,
Zador & Nicole
Robot Learning TAs